

Amogh Sahore 2547108 Machine Learning Lab 03

Data-cleaning assumptions

- Values between 0 and 1 in percentage columns are treated as fractions and multiplied by 100.
- Nonnumeric placeholders (for example, `Option 1`, `na`, and `-`) become missing values.
- GPA is expected on a 4-point scale. Values above 4 through 10 are converted from a 10-point scale; values above 10 through 100 are converted from a percentage scale.
- CIA values below 40 after normalization are considered implausible survey-entry errors and are excluded from the CIA experiment.
- Exact duplicate rows are removed. If a registration number occurs more than once, the latest response is retained.
- Missing values are handled per experiment by dropping only rows missing that experiment's X or Y values.

```
In [ ]: #step 1: Import libraries and configure reproducibility
from pathlib import Path
import pickle
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

RANDOM_STATE = 42
pd.set_option('display.max_columns', None)
```

```
In [ ]: #step 2: Load the Excel dataset and export the requested CSV dataset
DATA_FILE = Path('E:\\College\\MCA\\Trimester 4\\ML Labs\\LabProgs\\LAB03\\Department Awareness Survey (Responses).xlsx')
if not DATA_FILE.exists():
    raise FileNotFoundError(f'Place {DATA_FILE.name!r} in the same folder as this notebook')

raw_df = pd.read_excel(DATA_FILE)
raw_df.to_csv('student_survey.csv', index=False)
print('Loaded:', DATA_FILE)
print('CSV created: student_survey.csv')
```

Loaded: E:\College\MCA\Trimester 4\ML Labs\LabProgs\LAB03\Department Awareness Survey (Responses).xlsx

CSV created: student_survey.csv

```
In [7]: # Major step 3: Display the first five rows
raw_df.head()
```

Out[7]:

	Timestamp	Registration Number	Email	Job role that you are interested in\n	What is the minimum salary of students placed through campus (In LPA..respond as a number)
0	2026-06-22 08:49:31.362	2547123	jijaelza.jabi@mca.christuniversity.in	Software Engineer	4
1	2026-06-22 08:49:50.520	2547122	jinishaleema.rosario@mca.christuniversity.in	Data Engineer	7,00,000
2	2026-06-22 08:50:44.521	2547101	rajeev.chandar@mca.christuniversity.in	Software Engineer	6
3	2026-06-22 08:51:00.666	2547156	sounak.chakraborty@mca.christuniversity.in	Data Scientist	4.5
4	2026-06-22 08:51:42.357	2547148	samar.subhash@mca.christuniversity.in	Software Engineer	8

```
In [ ]: #step 4: Check dataset dimensions and column names
print(f'Rows: {raw_df.shape[0]}, Columns: {raw_df.shape[1]}')
for index, column in enumerate(raw_df.columns, start=1):
    print(f'{index:02d}. {column!r}')
```

Rows: 54, Columns: 15

01. 'Timestamp'
02. 'Registration Number '
03. 'Email'
04. 'Job role that you are interested in\n'
05. 'What is the minimum salary of students placed through campus (In LPA..respond as a number)'
06. 'What is the maximum salary of students placed through campus (In LPA..respond as a number)'
07. 'What is the median salary of students placed through campus (In LPA..respond as a number)'
08. 'Which is the highest paying company'
09. 'Rate your contribution towards extra curricular activities'
10. 'Rate your technical competencies'
11. 'What are your package expectations (LPA)'
12. 'your CIA % of last semester '
13. 'your GPA of last semester'
14. 'Your maximum attendance % till last semester'
15. 'Internships Interests'

```
In [ ]: #step 5: Identify missing values before cleaning
missing_before = raw_df.isna().sum().sort_values(ascending=False)
missing_before.to_frame('missing_count')
```

Out[]:

	missing_count
Timestamp	0
Registration Number	0
Email	0
Job role that you are interested in\n	0
What is the minimum salary of students placed through campus (In LPA..respond as a number)	0
What is the maximum salary of students placed through campus (In LPA..respond as a number)	0
What is the median salary of students placed through campus (In LPA..respond as a number)	0
Which is the highest paying company	0
Rate your contribution towards extra curricular activities	0
Rate your technical competencies	0
What are your package expectations (LPA)	0
your CIA % of last semester	0
your GPA of last semester	0
Your maximum attendance % till last semester	0
Internships Interests	0

```
In [10]: #step 6: Select and rename the required survey columns
def find_column(columns, required_terms):
    matches = [c for c in columns if all(term.lower() in c.lower() for term in required_terms)]
    if len(matches) != 1:
        raise ValueError(f'Expected one column containing {required_terms}, found: {matches}')
    return matches[0]

registration_col = find_column(raw_df.columns, ['registration'])
cia_col = find_column(raw_df.columns, ['cia', '%'])
gpa_col = find_column(raw_df.columns, ['gpa'])
attendance_col = find_column(raw_df.columns, ['attendance', '%'])
timestamp_col = find_column(raw_df.columns, ['timestamp'])

df = raw_df.rename(columns={
    registration_col: 'Registration_Number',
    cia_col: 'CIA_Percentage',
    gpa_col: 'GPA',
    attendance_col: 'Attendance_Percentage',
    timestamp_col: 'Timestamp',
}).copy()

df[['Timestamp', 'Registration_Number', 'CIA_Percentage', 'GPA', 'Attendance_Percentage']]
```

Out[10]:

	Timestamp	Registration_Number	CIA_Percentage	GPA	Attendance_Percentage
0	2026-06-22 08:49:31.362	2547123	78	3.24	Option 1
1	2026-06-22 08:49:50.520	2547122	64.78	3.2	Option 1
2	2026-06-22 08:50:44.521	2547101	80	3.6	Option 1
3	2026-06-22 08:51:00.666	2547156	70	3.2	92
4	2026-06-22 08:51:42.357	2547148	68	3.4	85

```
In [11]: #step 7: Remove exact duplicates and repeated registrations
rows_initial = len(df)
exact_duplicates = int(df.duplicated().sum())
df = df.drop_duplicates().copy()

df['Timestamp'] = pd.to_datetime(df['Timestamp'], errors='coerce')
df['Registration_Number'] = df['Registration_Number'].astype('string').str.strip()
df = df.sort_values('Timestamp').drop_duplicates(subset='Registration_Number', keep='last')

print('Initial rows:', rows_initial)
print('Exact duplicate rows removed:', exact_duplicates)
print('Repeated-registration rows removed:', rows_initial - exact_duplicates - len(df))
print('Rows remaining:', len(df))
```

Initial rows: 54

Exact duplicate rows removed: 0

Repeated-registration rows removed: 2

Rows remaining: 52

```
In [12]: #step 8: Convert the regression columns to numeric and normalize scales
numeric_columns = ['CIA_Percentage', 'GPA', 'Attendance_Percentage']
for column in numeric_columns:
    df[column] = pd.to_numeric(df[column], errors='coerce')

for column in ['CIA_Percentage', 'Attendance_Percentage']:
    fraction_mask = df[column].between(0, 1, inclusive='both')
    df.loc[fraction_mask, column] = df.loc[fraction_mask, column] * 100
    df.loc[~df[column].between(0, 100, inclusive='both'), column] = np.nan

gpa_from_10 = df['GPA'].gt(4) & df['GPA'].le(10)
gpa_from_100 = df['GPA'].gt(10) & df['GPA'].le(100)
df.loc[gpa_from_10, 'GPA'] = df.loc[gpa_from_10, 'GPA'] * (4 / 10)
df.loc[gpa_from_100, 'GPA'] = df.loc[gpa_from_100, 'GPA'] * (4 / 100)
df.loc[~df['GPA'].between(0, 4, inclusive='both'), 'GPA'] = np.nan

df[numeric_columns].dtypes
```

```
Out[12]: CIA_Percentage    float64
GPA                      float64
Attendance_Percentage    float64
dtype: object
```

```
In [13]: #step 9: Review null values after numeric conversion
df[numeric_columns].isna().sum().to_frame('missing_after_cleaning')
```

```
Out[13]:
```

missing_after_cleaning	
CIA_Percentage	1
GPA	1
Attendance_Percentage	4

```
In [14]: #step 10: Generate the statistical summary
df[numeric_columns].describe().T
```

```
Out[14]:
```

	count	mean	std	min	25%	50%	75%	max
CIA_Percentage	51.0	71.399608	15.228290	3.45	70.000	71.0	79.6650	89.0
GPA	51.0	3.492549	0.276383	2.90	3.385	3.5	3.6550	4.0
Attendance_Percentage	48.0	93.121458	4.807110	85.00	89.000	93.5	97.9925	100.0

```
In [15]: #step 11: Select independent and dependent variables for both experiments
cia_data = df.loc[df['CIA_Percentage'].ge(40), ['CIA_Percentage', 'GPA']].dropna().copy()
attendance_data = df[['Attendance_Percentage', 'GPA']].dropna().copy()

print('Experiment 1 - X: CIA_Percentage, Y: GPA, usable rows:', len(cia_data))
print('Experiment 2 - X: Attendance_Percentage, Y: GPA, usable rows:', len(attendance_data))
display(cia_data.head())
display(attendance_data.head())
```

Experiment 1 - X: CIA_Percentage, Y: GPA, usable rows: 48

Experiment 2 - X: Attendance_Percentage, Y: GPA, usable rows: 47

	CIA_Percentage	GPA
2	80.0	3.6
3	70.0	3.2
4	68.0	3.4
5	70.0	3.5
6	88.0	3.7

	Attendance_Percentage	GPA
3	92.0	3.2
4	85.0	3.4
5	87.0	3.5
7	90.0	4.0
8	89.0	3.4

```
In [16]: #step 12: Define the manual Ordinary Least Squares calculation
def manual_ols(x, y):
```

```

x = np.asarray(x, dtype=float).reshape(-1)
y = np.asarray(y, dtype=float).reshape(-1)
x_mean, y_mean = x.mean(), y.mean()
denominator = np.sum((x - x_mean) ** 2)
if np.isclose(denominator, 0):
    raise ValueError('OLS slope is undefined because all X values are identical.')
slope = np.sum((x - x_mean) * (y - y_mean)) / denominator
intercept = y_mean - slope * x_mean
return float(slope), float(intercept)

print('Slope formula:  $\Sigma((x - \bar{x})(y - \bar{y})) / \Sigma((x - \bar{x})^2)$ ')
print('Intercept formula:  $\bar{y} - \text{slope} \times \bar{x}$ ')

```

Slope formula: $\Sigma((x - \bar{x})(y - \bar{y})) / \Sigma((x - \bar{x})^2)$

Intercept formula: $\bar{y} - \text{slope} \times \bar{x}$

Experiment 1: CIA Percentage predicts GPA

```

In [17]: #step 13: Split CIA data into training and testing sets
X_cia = cia_data[['CIA_Percentage']]
y_cia = cia_data['GPA']
X_cia_train, X_cia_test, y_cia_train, y_cia_test = train_test_split(
    X_cia, y_cia, test_size=0.20, random_state=RANDOM_STATE
)
print('Training rows:', len(X_cia_train), ' | Testing rows:', len(X_cia_test))

```

Training rows: 38 | Testing rows: 10

```

In [18]: #step 14: Train Scikit-Learn regression for CIA → GPA
cia_model = LinearRegression()
cia_model.fit(X_cia_train, y_cia_train)
cia_sklearn_slope = float(cia_model.coef_[0])
cia_sklearn_intercept = float(cia_model.intercept_)
cia_sklearn_predictions = cia_model.predict(X_cia_test)
print(f'Slope: {cia_sklearn_slope:.10f}')
print(f'Intercept: {cia_sklearn_intercept:.10f}')
print(f'Equation: GPA = {cia_sklearn_intercept:.6f} + ({cia_sklearn_slope:.6f} × CIA_Percentage)')

```

Slope: 0.0110265790

Intercept: 2.6338863024

Equation: GPA = 2.633886 + (0.011027 × CIA_Percentage)

```

In [19]: #step 15: Compute the CIA regression equation manually using OLS
cia_ols_slope, cia_ols_intercept = manual_ols(X_cia_train.iloc[:, 0], y_cia_train)
cia_ols_predictions = cia_ols_intercept + cia_ols_slope * X_cia_test.iloc[:, 0].to_numpy()
print(f'Manual slope: {cia_ols_slope:.10f}')
print(f'Manual intercept: {cia_ols_intercept:.10f}')
print(f'Equation: GPA = {cia_ols_intercept:.6f} + ({cia_ols_slope:.6f} × CIA_Percentage)')

```

Manual slope: 0.0110265790

Manual intercept: 2.6338863024

Equation: GPA = 2.633886 + (0.011027 × CIA_Percentage)

```

In [20]: #step 16: Compare Scikit-Learn and manual OLS predictions for CIA
cia_comparison = pd.DataFrame({
    'CIA_Percentage': X_cia_test.iloc[:, 0].to_numpy(),
    'Actual_GPA': y_cia_test.to_numpy(),
    'Sklearn_Prediction': cia_sklearn_predictions,
    'Manual_OLS_Prediction': cia_ols_predictions,
})

```

```
cia_comparison['Absolute_Difference'] = (
    cia_comparison['Sklearn_Prediction'] - cia_comparison['Manual_OLS_Prediction']
).abs()
cia_comparison
```

Out[20]:

	CIA_Percentage	Actual_GPA	Sklearn_Prediction	Manual_OLS_Prediction	Absolute_Difference
0	80.0	3.86	3.516013	3.516013	0.0
1	77.0	3.77	3.482933	3.482933	0.0
2	77.0	3.60	3.482933	3.482933	0.0
3	78.0	3.74	3.493959	3.493959	0.0
4	85.0	3.60	3.571146	3.571146	0.0
5	80.0	3.80	3.516013	3.516013	0.0
6	89.0	3.67	3.615252	3.615252	0.0
7	80.0	3.64	3.516013	3.516013	0.0
8	88.0	3.70	3.604225	3.604225	0.0
9	70.0	3.40	3.405747	3.405747	0.0

In [21]:


```
#step 17: Evaluate the CIA model
cia_mse = mean_squared_error(y_cia_test, cia_sklearn_predictions)
cia_metrics = {
    'MAE': mean_absolute_error(y_cia_test, cia_sklearn_predictions),
    'MSE': cia_mse,
    'RMSE': cia_mse ** 0.5,
    'R2': r2_score(y_cia_test, cia_sklearn_predictions),
    'Maximum_method_difference': cia_comparison['Absolute_Difference'].max(),
}
pd.Series(cia_metrics, name='CIA → GPA')
```

Out[21]:

```
MAE          0.158726
MSE          0.038403
RMSE         0.195968
R2          -1.530528
Maximum_method_difference  0.000000
Name: CIA → GPA, dtype: float64
```

In [22]:

```
#step 18: Plot CIA percentage against GPA with the fitted line
x_line = np.linspace(cia_data['CIA_Percentage'].min(), cia_data['CIA_Percentage'].max(), 100)
plt.figure(figsize=(8, 5))
plt.scatter(cia_data['CIA_Percentage'], cia_data['GPA'], alpha=0.75, label='Cleaned observations')
plt.plot(x_line, cia_sklearn_intercept + cia_sklearn_slope * x_line, color='crimson', label='Fitted Line')
plt.xlabel('CIA Percentage')
plt.ylabel('GPA')
plt.title('Experiment 1: CIA Percentage vs GPA')
plt.legend()
plt.grid(alpha=0.2)
plt.show()
```



Experiment 2: Attendance Percentage predicts GPA

```
In [23]: #step 19: Split attendance data into training and testing sets
X_att = attendance_data[['Attendance_Percentage']]
y_att = attendance_data['GPA']
X_att_train, X_att_test, y_att_train, y_att_test = train_test_split(
    X_att, y_att, test_size=0.20, random_state=RANDOM_STATE
)
print('Training rows:', len(X_att_train), ' | Testing rows:', len(X_att_test))
```

Training rows: 37 | Testing rows: 10

```
In [24]: #step 20: Train Scikit-learn regression for attendance → GPA
attendance_model = LinearRegression()
attendance_model.fit(X_att_train, y_att_train)
attendance_sklearn_slope = float(attendance_model.coef_[0])
attendance_sklearn_intercept = float(attendance_model.intercept_)
attendance_sklearn_predictions = attendance_model.predict(X_att_test)
print(f'Slope: {attendance_sklearn_slope:.10f}')
print(f'Intercept: {attendance_sklearn_intercept:.10f}')
print(f'Equation: GPA = {attendance_sklearn_intercept:.6f} + ({attendance_sklearn_slope:.6f} × Attendance_Percentage)')
```

Slope: 0.0255145234

Intercept: 1.1030658120

Equation: GPA = 1.103066 + (0.025515 × Attendance_Percentage)

```
In [25]: #step 21: Compute the attendance regression equation manually using OLS
attendance_ols_slope, attendance_ols_intercept = manual_ols(X_att_train.iloc[:, 0], y_att_train)
attendance_ols_predictions = attendance_ols_intercept + attendance_ols_slope * X_att_test
print(f'Manual slope: {attendance_ols_slope:.10f}')
print(f'Manual intercept: {attendance_ols_intercept:.10f}')
print(f'Equation: GPA = {attendance_ols_intercept:.6f} + ({attendance_ols_slope:.6f} × Attendance_Percentage)')
```


Manual slope: 0.0255145234

Manual intercept: 1.1030658120

Equation: $GPA = 1.103066 + (0.025515 \times Attendance_Percentage)$

```
In [26]: #step 22: Compare Scikit-Learn and manual OLS predictions for attendance
attendance_comparison = pd.DataFrame({
    'Attendance_Percentage': X_att_test.iloc[:, 0].to_numpy(),
    'Actual_GPA': y_att_test.to_numpy(),
    'Sklearn_Prediction': attendance_sklearn_predictions,
    'Manual_OLS_Prediction': attendance_ols_predictions,
})
attendance_comparison['Absolute_Difference'] = (
    attendance_comparison['Sklearn_Prediction'] - attendance_comparison['Manual_OLS_Prediction']).abs()
attendance_comparison
```

```
Out[26]:
```

	Attendance_Percentage	Actual_GPA	Sklearn_Prediction	Manual_OLS_Prediction	Absolute_Diff
0	94.00	3.60	3.501431	3.501431	
1	88.00	3.50	3.348344	3.348344	
2	96.00	3.86	3.552460	3.552460	
3	98.71	3.60	3.621604	3.621604	
4	85.00	3.40	3.271800	3.271800	
5	89.00	3.30	3.373858	3.373858	
6	94.25	3.37	3.507810	3.507810	
7	85.00	3.60	3.271800	3.271800	
8	89.00	3.40	3.373858	3.373858	
9	97.00	3.60	3.577975	3.577975	


```
In [27]: #step 23: Evaluate the attendance model
attendance_mse = mean_squared_error(y_att_test, attendance_sklearn_predictions)
attendance_metrics = {
    'MAE': mean_absolute_error(y_att_test, attendance_sklearn_predictions),
    'MSE': attendance_mse,
    'RMSE': attendance_mse ** 0.5,
    'R2': r2_score(y_att_test, attendance_sklearn_predictions),
    'Maximum_method_difference': attendance_comparison['Absolute_Difference'].max(),
}
pd.Series(attendance_metrics, name='Attendance → GPA')
```

```
Out[27]:
```

MAE	0.129560
MSE	0.027753
RMSE	0.166592
R2	-0.150567
Maximum_method_difference	0.000000
Name: Attendance → GPA, dtype: float64	

```
In [28]: #step 24: Plot attendance percentage against GPA with the fitted line
x_line = np.linspace(attendance_data['Attendance_Percentage'].min(), attendance_data['Attendance_Percentage'].max(), 100)
plt.figure(figsize=(8, 5))
plt.scatter(attendance_data['Attendance_Percentage'], attendance_data['GPA'], alpha=0.75)
```

```
plt.plot(x_line, attendance_sklearn_intercept + attendance_sklearn_slope * x_line, color
plt.xlabel('Attendance Percentage')
plt.ylabel('GPA')
plt.title('Experiment 2: Attendance Percentage vs GPA')
plt.legend()
plt.grid(alpha=0.2)
plt.show()
```



Parameter saving and loading

```
In [29]: #step 25: Save both models' learned slope and intercept using Pickle
learned_parameters = {
    'CIA_to_GPA': {'slope': cia_sklearn_slope, 'intercept': cia_sklearn_intercept},
    'Attendance_to_GPA': {
        'slope': attendance_sklearn_slope,
        'intercept': attendance_sklearn_intercept,
    },
}

PICKLE_FILE = Path('linear_regression_weights.pkl')
with PICKLE_FILE.open('wb') as file:
    pickle.dump(learned_parameters, file)
print('Saved:', PICKLE_FILE)
learned_parameters
```

Saved: linear_regression_weights.pkl

```
Out[29]: {'CIA_to_GPA': {'slope': 0.011026579043724906,
    'intercept': 2.6338863023760415},
    'Attendance_to_GPA': {'slope': 0.02551452344903064,
    'intercept': 1.1030658119827317}}
```

```
In [30]: #step 26: Load the saved model parameters
with PICKLE_FILE.open('rb') as file:
    loaded_parameters = pickle.load(file)
loaded_parameters
```

```
Out[30]: {'CIA_to_GPA': {'slope': 0.011026579043724906,
    'intercept': 2.6338863023760415},
    'Attendance_to_GPA': {'slope': 0.02551452344903064,
    'intercept': 1.1030658119827317}}
```

```
In [31]: #step 27: Use Loaded parameters for new predictions
def predict_with_saved_parameters(experiment_name, x_value):
    parameters = loaded_parameters[experiment_name]
    return parameters['intercept'] + parameters['slope'] * x_value

sample_cia = 75.0
sample_attendance = 90.0
print(f'Predicted GPA for CIA {sample_cia:.1f}%: '
      f'{predict_with_saved_parameters("CIA_to_GPA", sample_cia):.4f}')
print(f'Predicted GPA for attendance {sample_attendance:.1f}%: '
      f'{predict_with_saved_parameters("Attendance_to_GPA", sample_attendance):.4f}')
```

Predicted GPA for CIA 75.0%: 3.4609

Predicted GPA for attendance 90.0%: 3.3994

```
In [32]: #step 28: Produce final observations and inference
method_tolerance = 1e-10
results = pd.DataFrame([
    {
        'Experiment': 'CIA → GPA',
        'Slope': cia_sklearn_slope,
        'Intercept': cia_sklearn_intercept,
        **cia_metrics,
    },
    {
        'Experiment': 'Attendance → GPA',
        'Slope': attendance_sklearn_slope,
        'Intercept': attendance_sklearn_intercept,
        **attendance_metrics,
    },
])
display(results)

for _, row in results.iterrows():
    same = row['Maximum_method_difference'] < method_tolerance
    direction = 'positive' if row['Slope'] > 0 else 'negative' if row['Slope'] < 0 else
    print(f"{row['Experiment']}: Scikit-learn and manual OLS predictions "
          f"{ 'match to floating-point precision' if same else 'need investigation' }.")
    print(f" The fitted line has a {direction} slope ({row['Slope']:.6f}) "
          f"and test R² = {row['R2']:.4f}.")

best = results.loc[results['R2'].idxmax()]
print(f"Among these two one-variable models, {best['Experiment']} has the higher test R². "
      "Interpret cautiously because the dataset is small and self-reported; association
```

	Experiment	Slope	Intercept	MAE	MSE	RMSE	R2	Maximum_method_dif
0	CIA → GPA	0.011027	2.633886	0.158726	0.038403	0.195968	-1.530528	
1	Attendance → GPA	0.025515	1.103066	0.129560	0.027753	0.166592	-0.150567	

◀  ▶

CIA → GPA: Scikit-learn and manual OLS predictions match to floating-point precision.
 The fitted line has a positive slope (0.011027) and test $R^2 = -1.5305$.
 Attendance → GPA: Scikit-learn and manual OLS predictions match to floating-point precision.
 The fitted line has a positive slope (0.025515) and test $R^2 = -0.1506$.
 Among these two one-variable models, Attendance → GPA has the higher test R^2 . Interpret cautiously because the dataset is small and self-reported; association does not establish causation.